

PM INTERN ASSIGNMENT for SKILLOVILLA

Submitted by Ravi Vinayak (vinayakravi01@gmail.com)

Ditto Insurance (joinditto.in)

Sections A, B, C & D — Sales Funnel Research · Platform Comparison · AI Caller Design · Execution Plan

Section – A : SALES FUNNEL RESEARCH


QUESTION 1

Map the Sales Funnel & Identify the AI Calling Opportunity

How Ditto's product actually works

Research basis — what Ditto actually does: The primary CTA on <https://joinditto.in/> is "Book a Free Call," not "Compare Quotes" — this shifts the user from transaction to consultation mode by design. Ditto has 800,000+ customers advised, 25,000+ booked consultations/month, and a 4.9/5 Google rating from 22,000+ reviews. <https://joinditto.in/>. https://founderpin.com/startup_story/ditto/ Advisors are IRDAI-certified. Post-sale, Ditto's claims concierge handles 500+ claim settlements monthly — this is both a service and a retention mechanism. Acquisition is via Finshot's newsletter (1M+ subscribers), YouTube, SEO blogs, Google Ads (high-intent), and Meta Ads targeting 25–40 age group in Tier 1 cities.

Ditto's Complete Sales Funnel



Conversion rates are estimates derived from industry benchmarks and Ditto's publicly available scale data. No internal Ditto funnel data is public.

Stages	What Actually Happens at Ditto	Drop-off Risk	AI-Calling fit ?
1. Visitor → Lead	User arrives via SEO, Finshots, YouTube, or ads. Sees "Book a Free Call" CTA. Fills form with name, phone, insurance type.	Medium Finance/insurance landing pages convert at 2–10% depending on intent; paid search (high-intent) converts at ~10% median. https://unbounce.com/conversion-benchmark-report/	No — top of funnel, not calling-appropriate
2. Lead → Qualified	System (or first advisor call) checks product fit, urgency, and contactability. Ditto's job postings confirm advisors handle "inbound, not outbound" — qualification happens at or before the first call.	Medium Generic insurance SaaS benchmarks: 28% MQL→SQL conversion (mid-funnel is the biggest bottleneck). https://firstpagesage.com/seo-blog/saas-marketing-channels-compared/	Partial – AI can prequalify and flag low intent lags
3. Booked → First Contact	User submits form but advisor hasn't spoken to them yet. Speed-to-contact is critical — research shows 85% of leads don't convert if not reached within 5 minutes. At Ditto's stated 25,000+ calls/month scale, inbound lead volume likely outstrips manual follow-up capacity during peak periods.	HIGH — primary bottleneck The assignment brief explicitly confirms: "many inbound leads never get connected to an advisor on time"	YES — AI Priority #1
4. Booked → Attended	User booked a slot but doesn't pick up or misses the call. Advisors confirmed (https://www.glassdoor.co.in/ditto) that clients call back at off-hours, suggesting scheduling compliance is inconsistent. Glassdoor's Ditto review The assignment confirms "a significant number of scheduled consultations are missed."	HIGH — confirmed in brief No-show rates in high-consideration service funnels typically 25–40%	YES — AI Priority #2
5. Consultation → Recommendation	IRDAI-certified advisor spends 20–40 min understanding medical history, family structure, and budget. Personalized plan comparison. This is Ditto's core product — the "trust" moment. (https://irdai.gov.in/)	Low Users who make it here are high-intent and already trust Ditto's brand	No — human judgment required; AI here would damage brand
6. Recommendation → Application	User shares personal details, fills proposer form, uploads documents. Form friction and paperwork delay are the main risk here. Some advisor handholding required.	Medium Regulatory document requirements cause drop-off in Indian insurance applications	Partial — AI reminder/document nudge possible (lower priority)
7. Application → Issuance	Underwriting, medicals, payment, insurer approval. Operationally complex; insurer-dependent. Ditto advisors help chase these. joinditto.in/contact-us	Medium Insurer and underwriting timelines outside Ditto's control	No — human coordination with insurers needed
8. Issuance → Renewal / Claims	Ditto's claims concierge: 500+ claim settlements/month. Dedicated round-the-clock support team. (https://joinditto.in/) in homepage This is a retention + NPS driver and a key differentiator vs Policybazaar.	Low High-touch support reduces churn and drives referrals	Partial — AI reminder for renewal nudges (future phase)

AI Calling Priority: Stages to Target

P1

Lead → First Contact (Speed-to-Call)

The highest-volume, most time-sensitive, and most fixable leakage point. Ditto's brief explicitly names this as a problem: "many inbound leads never get connected to an advisor on time." AI can call within seconds of form submission, qualify intent, and book a human consultation — without giving any insurance advice. This is structurally safe for Ditto's brand because the AI's job is narrow: connect and schedule, nothing more. The task is repetitive, scalable, and low-risk.

P2

Consultation No-Show Recovery

The second confirmed problem in the brief: "a significant number of scheduled consultations are missed." An AI caller that rings the user post-no-show, offers an immediate reschedule, and sends a WhatsApp confirmation handles this at zero marginal advisor cost. Like P1, the task is structured, bounded, and doesn't require insurance knowledge — just intent re-confirmation and slot rebooking.

P3

Post-Consultation Follow-Up (Phase 2 — Document Reminders, Application Nudges)

Lower priority because it sits deeper in the funnel where users are already converted and the advisor relationship is established. Recommended as Phase 2 after P1 and P2 are validated. Narrower pilot scope means faster learning and lower brand-risk exposure.

Why NOT the consultation itself? Stages 5–6 (recommendation and application) are where Ditto's human advisors spend 20–40 minutes building trust, personalizing advice, and answering complex policy questions. AI calling here would violate Ditto's core brand positioning — IRDAI regulations also require licensed advisor involvement for insurance recommendation. The original answer correctly identified this boundary. It is preserved here.

Section – B :

Platform Research & Comparison

QUESTION 2

Research, Compare & Recommend AI Calling Platforms

- **Retell AI:** Pricing page verified — starts at **\$0.07+/min** for AI Voice Agents (pay-as-you-go). Total all-in cost (voice + LLM + telephony) typically ranges from **\$0.13–\$0.31/min**. Enterprise plans (above \$3K/month) reduce the base cost to approximately **\$0.05/min**.

Sources: <https://www.retellai.com/pricing> ;

- **Vapi:** Advertised base pricing is **\$0.05/min** (hosting only). Mandatory add-ons — including telephony via Twilio (~\$0.008–\$0.014/min), speech-to-text (~\$0.01/min), text-to-speech (~\$0.02/min), and LLM usage (~\$0.05–\$0.08/min) — increase the actual total cost to approximately **\$0.18–\$0.33/min**.

Sources: <https://www.cloudtalk.io/blog/vapi-ai-pricing/> ; <https://telnyx.com/resources/vapi-pricing>

- **Plivo:** India telephony voice infrastructure pricing starts at **\$0.0040/min**. Plivo also offers an AI Agent product with pay-as-you-go billing and **\$10 free credits**. The platform operates across **7 global PoPs including India** and provides a **99.99% uptime SLA**. Full AI agent pricing additionally requires LLM + STT/TTS service costs.

Sources: <https://www.plivo.com/voice/pricing/in/> ; <https://www.plivo.com/blog/ai-voice-agents/>

PLATFORM	VERIFIED PRICING	KEY FEATURES	LANGUAGE / VOICE	CRM & INTEGRATION	LIMITATIONS FOR DITTO
Retell AI RECOMMENDED	\$0.07+/min All-in (LLM + telephony): \$0.13–\$0.31/min. Enterprise ~\$0.05/min at >\$3K/month. \$10 free credits to start (~60 pilot calls). No platform fee.	PAYG, no forced bundles. Structured conversation flow control (critical for Dia's bounded script). Visual call builder (2025). Low-latency streaming. Interruption handling. Batch outbound calling. SOC 2 + HIPAA compliant. 30+ language support.	30+ languages. Hindi via ElevenLabs/Plivo HT. Manual config for Hinglish. Indian accent STT quality depends on provider selection.	Developer-first; no native CRM — integrates via webhooks. Compatible with HubSpot, Leadsquared via custom webhook. WhatsApp trigger via Zapier/n8n.	Non-technical teams face a learning curve. Occasional latency spikes under load. India telephony quality needs testing — route via Twilio or Plivo SIP. Billing complexity at scale.
Vapi ALTERNATIVE	\$0.05/min (base only) △Real all-in cost: \$0.18–\$0.33/min. Headline excludes telephony, STT, TTS, LLM — all mandatory. Plans from \$500/month. Enterprise \$40K–\$70K/year.	Highly flexible API-first. Bring-your-own LLM, STT, TTS. Sub-600ms response time. Strong developer ecosystem. Good for custom agent workflows. Concurrency controls.	Flexible but manual. Mix-and-match voice stack. Good for English + Hindi but requires manual config per language.	No native CRM; developer-managed webhooks. Integrates with Twilio, Vonage, Telnyx for telephony.	Pricing opacity is the biggest risk — \$0.05/min headline causes budget overruns. Trustpilot 2.6 reflects mixed billing/support experiences(. https://telnyx.com/resources/vapi-pricing) Requires strong engineering bandwidth. Not ideal for quick pilot.
Plivo INFRASTRUCTURE	\$0.0040/min Telephony infrastructure (India). AI Agent product PAYG with \$10 free credits. Plans from \$25/month base + usage. Full all-in requires stacking LLM + STT/TTS.	Proven telephony infrastructure. 7 global PoPs, 99.99% uptime. Strong India coverage with local number support. Full AI agent product (no-code + low-code). OpenAI + Deepgram + ElevenLabs integrations. Outbound + inbound support.	Global + India. India-specific local numbers. STT via Deepgram (strong for Indian accents). ElevenLabs TTS.	No native CRM — developer-managed API. WhatsApp Business API also available (useful for Ditto's post-call confirmation flow from a single platform).	More of a communications infrastructure layer than a ready-made AI agent. More engineering setup than Retell to go live. Better suited for scale phase after pilot is validated.

Recommendation for Ditto's Pilot

Start with Retell AI. For a pilot optimizing for speed-to-launch, predictable per-minute economics, structured conversation control and no surprise billing, Retell is the strongest fit. The \$10 free credits (~60 pilot calls), visual flow builder, and PAYG pricing let a PM-led team get to a live test in days, not weeks. Structured flow agents are specifically built for bounded scripts like Dia's.

Vapi as fallback if Retell's India telephony quality underperforms after testing. But budget at \$0.18–\$0.33/min from day one — never use the \$0.05/min headline in ROI calculations.

Plivo for the scale phase. Once the pilot validates the script and trust-preservation, Plivo's \$0.004/min telephony + AI agent layer becomes the most cost-efficient stack at volume. It also handles both voice calling and WhatsApp confirmation from a single platform.

Optimizing for: speed-to-pilot · transparent pilot cost · brand-safe structured call control · India telephony reliability.

Designing the AI caller

QUESTION 3

AI Caller Persona & Call Scripts

Dia

DITTO'S AI SCHEDULING ADVISOR

Dia is a warm, unhurried presence on the phone — the opposite of the typical insurance cold-caller. She doesn't push products, doesn't read from a script mechanically, and never creates urgency pressure. She sounds like someone who genuinely wants to help you get to the right expert, not close a deal.

Tone

Calm, clear, friendly — never rushed

Pace

Slightly slower than average; deliberate pauses

Language

Simple English; Hindi fallback available

Identity

Always identifies as AI upfront

Brand cue

No spam, no push — scheduling only

Escalation

Hands off to human advisor on first serious question

Target stage: Lead-to-First-Connect — an inbound lead submitted their name & phone number on Ditto's website but has not yet spoken with any advisor. Dia's job is to confirm intent, handle basic qualification, and lock in a consultation slot. She does not give insurance advice.

CALL SCRIPT — LEAD-TO-FIRST-CONNECT ⏻ ~50 sec speaking	
DIA	"Hi, is this [First Name]? I'm Dia, an AI assistant from Ditto Insurance. You'd recently asked us about health or term insurance — I'm calling to help you book a free call with one of our human advisors. This will take just a minute. Is now an okay time to talk?"
USER	"Yes, go ahead." / "Who is this again?"
DIA	"Great! Just to make sure we connect you with the right advisor — are you looking at health insurance, term life insurance, or are you still exploring both?"
USER	"Health insurance." / "Both, actually."
DIA	"Perfect. Our advisors are IRDAI-certified and this is completely free — they won't push you to buy anything. I can book you a slot for tomorrow morning at 10 AM or 11 AM. Which works better for you?"
USER	"11 is fine." / "Can I get an afternoon slot?"
DIA	"Done! I've booked you for [confirmed slot]. Your advisor will call you on this number. You'll also get a WhatsApp confirmation in a few minutes. Is there anything else you'd like me to note for them?"
USER	"No, that's fine. Thanks."
DIA	"Great, looking forward to it. Take care, [First Name]!"
💡 Design note: Dia never mentions policy names, premiums, or gives any insurance advice. If asked, she says: "That's exactly what our advisor will help you with — I'll make sure they're ready." Policy-related questions are a hard escalation trigger.	

Brand alignment check: Ditto's promise is "no spam, no push." Dia enforces this structurally — she has no access to policy data, cannot recommend or quote products, and terminates politely if the user opts out at any point. The word "free" appears once, deliberately, because it's a real differentiator — not a sales hook.

QUESTION 4

Edge cases – What could go Wrong

Edge cases are organized by category. Each row shows the trigger condition, Dia's handling flow, and whether the outcome is automated or requires human escalation.

1. Customer Availability and Timing

TRIGGER	HANDLING FLOW
"I'm busy right now"	Dia says "No problem at all — when is a better time for me to call back?" and offers 2–3 specific time options. Schedules a callback, confirms via WhatsApp.
No answer / rings out	System waits 24 hours, retries once (max 2 attempts total). If still no answer, falls to WhatsApp message: "Hi [Name], we tried calling — here's a link to book your consultation at your convenience."
Call drops mid-conversation	System detects disconnection, auto-sends WhatsApp: "Hi [Name], our call got cut off. Here's your booking link: [link]." Does not retry the call automatically (avoids feeling aggressive).
User is in a noisy environment	If STT confidence falls below threshold for 2+ consecutive turns, Dia says: "The line seems a bit noisy — shall I send you a booking link on WhatsApp instead so it's easier?"

2. Conversational Dead Ends

TRIGGER	HANDLING FLOW
Silence / no response for 4+ seconds	Dia gently prompts once: "Are you there? I just wanted to confirm your booking." If silence continues, says goodbye warmly and terminates call. Logs as "non-responsive" — system sends WhatsApp.
User says only "Hm?", "What?", "Hello?" repeatedly	After 2 failed clarification attempts, Dia offers to call back: "It seems we may have a bad connection — I'll try to reach you again. Take care!" Logs as audio quality failure.
Unrelated conversation topic (e.g., user thinks it's a different call)	Dia re-introduces herself clearly, confirms name, and asks if they're interested in insurance guidance. If user says no or is confused, ends call politely: "No worries! I'll let you go."
User asks same question repeatedly (looping)	After 2 loops, Dia says: "I want to make sure you get the right answer — let me have one of our advisors call you back in a few minutes." Triggers immediate human callback flag.

3. Language Switches

TRIGGER	HANDLING FLOW
User speaks in Hindi	Dia switches to Hindi mid-call using pre-loaded Hindi script variant. Core booking flow mirrors the English script.
User mixes Hindi and English (Hinglish)	Dia continues in Hinglish naturally. The agent is designed to detect code-switching and respond in kind.
User speaks in a regional language (Tamil, Kannada, Marathi, etc.)	Dia says in the detected language (using basic phrase detection): "I'll arrange for someone who speaks [language] to call you. Is that okay?" Routes to language-matched advisor queue.

4. Escalation Triggers — Policy & Advice Questions

TRIGGER	HANDLING FLOW
Asks which insurer or plan to buy	Dia says: "That's a great question for your advisor — they'll compare plans based on your age, city, and needs. Should I book that slot now?" Redirects to booking.
Asks about premium amounts or coverage limits	Same as above. Dia never quotes figures. Hard boundary. If pressed repeatedly, flags for human callback.
Mentions existing policy problem or claim issue	Dia immediately says: "That sounds important — I'm connecting you to Ditto's claims team right now." Flags as urgent, routes to claims support queue in CRM.
User is clearly distressed or upset	Dia slows down, says: "I understand — I want to make sure you get the right person to help. Can I have someone from our team call you today?" Flags as priority human callback.

5. DND, Opt-Out & TRAI Compliance

TRIGGER	HANDLING FLOW
User says "Don't call me again" / "Remove me"	Dia says: "Absolutely, I've noted that. You won't receive any more calls from us. Take care!" System immediately marks lead as DNC (Do Not Call) in CRM. No further outreach.
Number is on TRAI DND registry (detected pre-call)	Pre-call screening checks DND status via TRAI API. Call is not placed. System sends an opt-in WhatsApp message instead: "You can book a consultation here whenever you're ready: [link]."
User asks "Is this a recorded call?"	Dia says: "Yes, this call may be recorded for quality and compliance. Is that okay?" If user says no, call ends, lead is flagged for WhatsApp-only outreach.

6. Voicemail and IVR systems

TRIGGER	HANDLING FLOW
Call goes to voicemail	Dia detects voicemail via AMD (Answering Machine Detection). Leaves a short message: "Hi [Name], this is Dia from Ditto Insurance. You'd asked about insurance — I'll send you a WhatsApp to book at your convenience. Take care!" Then immediately sends WhatsApp.
Call is answered by another person (family member, secretary)	Dia asks: "Is [First Name] available? I'm Dia from Ditto Insurance — they'd asked us about insurance guidance." If person says they're unavailable, offers callback time collection.
Call reaches an IVR / automated system	AMD detects non-human response, terminates call, logs as "IVR/bad number," flags for manual review.

7. Technical Failures

TRIGGER	HANDLING FLOW
STT confidence consistently low (bad audio quality)	After 2 low-confidence turns, Dia offers WhatsApp alternative. If user accepts, sends booking link. Logs call as "audio failure."
LLM response latency exceeds 3 seconds	System plays a soft filler: "Just a moment..." If latency persists, Dia tells user she'll send a WhatsApp with scheduling link. Flags for infra review.
Calendar booking API fails mid-call	Dia says: "I'm having a small technical hiccup — I'll have someone confirm your slot over WhatsApp in the next 10 minutes." System sends human advisor a notification to manually book.

Design principle behind all edge cases: Every failure path ends in either a WhatsApp booking link or a human advisor callback — never in a dead end. Dia's job is to be the bridge, not the bottleneck. This is especially critical for Ditto, where trust erosion from a bad AI call would undo the brand's most important differentiator.

Section – C :

Execution Plan

QUESTION 5

Step by Step Execution plan

1. Internal Alignment & Scoping

Week 1

Document the problem with data: current lead-to-contact rate, % of no-shows, advisor hours spent on follow-ups (source: CRM logs, advisor interviews). Define pilot's exact scope — stage, volume, geographies, languages.

- Confirm with founders: what's the hard "no" for AI behavior? (e.g., never quote premiums, always identify as AI)
- Define success thresholds: minimum acceptable booking rate, maximum acceptable complaint rate
- Get legal/compliance sign-off on TRAI regulations for AI outbound calls in India

2. Platform Vendor Selection & Contracting

Week 1–2

Issue RFQ to Retell AI, Vapi, and Plivo. Evaluate on: India telephony quality, Hindi/Hinglish support, CRM (Leadsquared or HubSpot) integration, PAYG pricing (no minimums), TRAI outbound compliance.

- Request sandbox access from all 3 vendors within 5 business days
- Run internal test calls — evaluate voice naturalness, latency, STT accuracy for Indian accents
- Recommendation: **Retell AI as primary** (transparent PAYG at ~\$0.07/min, structured flow control, fastest time-to-live); Vapi as fallback; Plivo for scale phase
- Sign contract with data processing addendum covering call recording consent and DPDPA 2023 compliance

3. Script Finalization & Persona QA

Week 2

Finalize Dia's script (English + Hindi variants) reviewed by: 3 Ditto advisors (for brand tone fit), 1 IRDAI compliance reviewer (for regulatory language), 5 target users in a mock call test.

- Record 20 test calls using sandbox and review transcripts for: tone, pacing, escalation accuracy, opt-out handling
- Score each call against a rubric: did Dia correctly escalate policy questions? Did opt-out work? Did she identify as AI?
- Iterate until 95%+ of test calls pass the rubric

4. Integration Build

Week 2–4

Engineering builds the integration layer: Lead CRM → AI Caller Platform → Advisor Calendar → WhatsApp API → Call Logging DB.

- Input trigger: new lead marked "uncontacted" in CRM after 15 min → auto-dispatches Dia call (within calling hours only)
- Calendar API integration: Dia reads advisor availability and books slots in real-time
- WhatsApp API (Interakt or Wati): sends confirmation/fallback messages post-call
- Call log database: captures all fields defined in Q6 output format
- DNC/DND pre-screening: TRAI DND API check before every call dispatch
- Build internal ops dashboard: live call monitoring, flagged calls, escalation queue

5. Internal Beta — Dogfooding

Week 4

Before any real leads, run Dia on internal team members and existing customers who opt in voluntarily. Goal: catch integration bugs, test edge cases, validate escalation flows.

- 50 internal test calls across different call scenarios
- Manual review of all 50 transcripts by PM + 2 advisors
- Fix all critical bugs before any live lead touches the system

6. Controlled Pilot — Cohort A (A/B Test)

Week 5–6

Run Dia on 10–15% of new inbound leads (English-speaking, metro cities — Delhi, Bengaluru, Mumbai, Hyderabad). The remaining 85–90% continue with manual advisor follow-up as the control group. True A/B test.

- Target: 200–300 calls in 2 weeks
- Monitor daily: booking rate, complaint rate, escalation rate, call duration
- Advisor team reviews 20% of Dia call transcripts daily for brand compliance
- Calling hours: strictly 9 AM – 7 PM; no calls on Sundays

7. Review, Iterate & Decision Gate

Week 7

Pull pilot metrics. Compare AI cohort vs control: booking rate, show rate, policy conversion rate, complaint rate, NPS. Present findings to founders with clear go/no-go recommendation.

- If booking rate ≥ 30% higher than control AND complaints < 1%: expand to 50% of leads
- If booking rate is flat OR complaints > 1%: pause, root-cause, re-script
- If escalation rate is high: brief advisor team to handle increased AI-sourced call volume

8. Expand & Optimize (Hindi + No-Show Recovery)

Week 8–12

If pilot passes gate: activate Hindi script, expand to Tier-2 cities, add no-show recovery use case. Run A/B tests on call timing, opening lines, and slot-offer phrasing.

- Continuous: weekly transcript sampling, monthly KPI review with founders
- Quarterly: revisit vendor pricing, evaluate whether to consolidate platforms

Calling Time Strategy

WINDOW	TIME	RATIONALE	PRIORITY
Morning prime	9:00 – 11:30 AM	Users are settled but not in deep work. Highest pick-up rate for working professionals.	High
Lunch window	12:30 – 1:30 PM	Short availability window; works for salaried workers on lunch break.	Medium
Evening	5:30 – 7:00 PM	Post-work availability for salaried professionals. 2nd highest pick-up window in Indian insurance outreach.	High
Avoid	Before 9 AM / After 7 PM / Sundays	TRAI regulations + brand trust. Calling outside these hours is intrusive and directly contradicts Ditto's no-spam positioning.	Never

QUESTION 6

Input and Output Data Formats

INPUT **Lead Data Consumed by AI System** Source: CRM (Leadsquared / HubSpot) → webhook trigger → AI Caller platform. Format: JSON, sent per-lead on creation or status change to "uncontacted".

```
{
  // Lead Identity
  "lead_id": "DTT-2024-089234",
  "name": "Priya Sharma",
```

```

"phone": "+919876543210",
"email": "priya.s@email.com",

// Lead Context (from form submission)
"insurance_interest": "health", // "health" | "term" | "both" | "unknown"
"city": "Bengaluru",
"preferred_language": "en", // "en" | "hi" | "unknown"
"lead_source": "organic_seo", // "seo" | "youtube" | "referral" | "social"
"submitted_at": "2024-01-15T09:42:00+05:30",

// Compliance & Calling Rules
"dnd_status": false, // pre-checked against TRAI DND API
"consent_given": true, // form opt-in confirmation
"call_attempts_remaining": 2, // max 2 AI call attempts per lead
"preferred_callback_time": null, // if user specified, else null

// Advisor Availability (fetched from calendar API)
"available_slots": [
  "2024-01-16T10:00:00+05:30",
  "2024-01-16T11:00:00+05:30",
  "2024-01-16T17:30:00+05:30"
]
}

```

OUTPUT **Call Result Logged After Each Interaction**

Written to call log DB and synced back to CRM as lead activity. Visible to advisors on the sales dashboard. Format: JSON.

```

{
  // Call Identity
  "call_id": "CALL-20240115-089234-01",
  "lead_id": "DTT-2024-089234",
  "called_at": "2024-01-15T10:05:22+05:30",
  "call_duration_seconds": 68,

  // Outcome
  "call_outcome": "booked",
  // Options: "booked" | "callback_requested" | "opted_out" |
  // "no_answer" | "voicemail" | "wrong_number" |
  // "escalated_human" | "technical_failure" | "busy"

  "booked_slot": "2024-01-16T11:00:00+05:30", // null if not booked
  "assigned_advisor_id": "ADV-042",

  // Conversation Intelligence
  "insurance_type_confirmed": "health",
  "language_used": "en",
  "escalation_triggered": false,
  "escalation_reason": null,
  // e.g., "asked_for_policy_advice" | "distressed_user" |
  // "claims_query" | "technical_failure"

  // Trust & Compliance Signals
  "opted_out": false,
  "user_asked_if_ai": false, // flag for analysis: is disclosure friction rising?
  "complaint_flagged": false,
  "recording_url": "https://storage.ditto.in/calls/...",
  "transcript_id": "TXN-20240115-089234",

  // Follow-up Actions Triggered
  "whatsapp_confirmation_sent": true,
  "next_action": "advisor_call_scheduled",
  "notes_for_advisor": "User interested in health insurance for family of 4. Mentioned pre-existing BP condition."
}

```

What advisors see on the dashboard: A clean card per lead showing — Name, booked slot, insurance type, language, any escalation flag, and Dia's notes for the advisor. One click plays the call recording. Advisors can mark the lead as "ready for consultation" with zero manual data entry.

QUESTION 7

KPIs – Measuring Success Ongoing

KPI 1

Lead-to-Consultation Booking Rate

This is the primary north star for Dia's existence. It measures what fraction of inbound leads who receive a Dia call successfully book a consultation. Compare against the control cohort (manual follow-up) weekly to measure true incremental lift.

Target: $\geq 35\%$ booking rate (vs $\sim 20\%$ baseline for manual)

KPI 2

Consultation Show Rate (Post-AI Booking)

A booking means nothing if the user ghosts the advisor. This KPI tracks whether AI-booked consultations are actually attended. A low show rate signals that Dia is booking slots without true intent confirmation — a script and persona problem, not just a timing problem.

Target: $\geq 65\%$ show rate (AI-booked consultations)

KPI 3

Complaint & Opt-Out Rate

Ditto's #1 brand asset is trust. This KPI tracks the fraction of Dia calls that result in a DNC request, a negative feedback submission, or a formal complaint. Any spike above 1% is an immediate pause signal. Tracked weekly, reviewed by founders monthly.

Target: $< 1\%$ complaints; $< 8\%$ opt-outs per call cohort

KPI 4

Lead-to-Consultation Booking Rate

This is the primary north star for Dia's existence. It measures what fraction of inbound leads who receive a Dia call successfully book a consultation. Compare against the control cohort (manual follow-up) weekly to measure true incremental lift.

Target: $\geq 35\%$ booking rate (vs $\sim 20\%$ baseline for manual follow-up)

KPI 5



Downstream Policy Conversion Rate (AI-Sourced vs Control)

The ultimate proof that AI-assisted lead connect doesn't just improve bookings but flows through to actual policy purchases. Tracks what % of leads who were first connected by Dia eventually purchase a policy, compared to leads followed up manually. This metric takes 60–90 days to materialize but is the strongest signal of long-term ROI. It also surfaces whether Dia's qualification accuracy (health vs term, intent level) translates to real buyer quality downstream.

Target: Policy conversion parity or better vs manual follow-up cohort at 90 days

Reporting cadence: KPIs 1–3 reviewed weekly (PM + Head of Sales). KPIs 4–5 reviewed monthly (PM + CEO). All five KPIs on a shared dashboard accessible to founders, advisors and engineering in real-time.

QUESTION 8

Sales Manager Dashboard**Dashboard video link :**

https://drive.google.com/file/d/1sSG6qagAuPATII_wFNKlfp4DHrk9zIp9/view?usp=sharing

Dashboard video walkthrough

QUESTION 9

Design Decisions — Why Each Section Exists

TOP STRIP — 5 KPIS

Headline Numbers in 10 Seconds

The five metrics chosen — Consultations Booked, Lead→Contact Rate, No-Show Rate, Advisor Hours Saved, Post-AI CSAT — map directly to the 5 KPIs proposed in Section D (Q7). The manager doesn't need to open a report; the daily number is the headline. Color coding (green/amber/red) enables instant status reads without reading numbers.

ALERT STRIP

Surfaced Before Anything Else

The alert strip sits above the fold and fires when a KPI crosses its threshold. Critically, it also shows what auto-action was taken (e.g., "Recovery campaign auto-triggered for 47 leads") — so the manager knows whether to intervene or just monitor. The strip is dismissible but logged.

TRUST SIGNALS PANEL

Treated as a First-Class Metric

Most AI calling dashboards only track conversion. For Ditto this would be wrong — brand erosion is existential. The trust panel tracks: CSAT after AI contact, opt-out rate, "Is this AI?" question rate, disclosure acceptance rate, complaint flags, and — critically — whether the AI ever crossed its script boundary into policy advice. Zero boundary violations is a daily non-negotiable.

ADVISOR QUEUE

The Handoff Is the Product

The dashboard closes the loop between AI and human. Dia's job ends when the advisor's begins. The queue panel shows each advisor's upcoming AI-booked load and flags overloaded advisors in real time. The "redistribute" action prevents Ditto's most critical experience — the human consultation — from degrading because one advisor got 18 calls queued at once.

CALLING WINDOW HEATMAP

Data-Driven Schedule Optimization

The 7-day heatmap shows booking rate by hour of day. The manager uses this to adjust AI calling schedules weekly — concentrating calls in the 10–11 AM and 4–5 PM peak windows, and pausing or reducing calls during 12–2 PM when connection rates historically drop. This is a strategic tool, not just monitoring.

LIVE CALL LOG

Operational Transparency

The manager can see the last 12 calls in real time: lead name, call type (first-connect vs no-show recovery), duration, outcome, and AI confidence score. The confidence score is especially important — a cluster of low-confidence calls may signal a script edge case the team needs to review. Escalated calls are flagged for transcript review.

Full Trade-Off Register — What Was Chosen and Why

DECISION	CHOSE	REASONING & WHAT WAS TRADED OFF
AI scope	NOT IN CONSULTATION	An AI that recommends policies would save more advisor time per lead — but it would require IRDAI licensing, carry mis-selling risk and erode the trust brand that is Ditto's only moat. The trade-off: lower efficiency gain in exchange for zero brand risk and zero regulatory exposure.
Platform	RETELL OVER VAPI	Vapi is more flexible and beloved by developers. Retell is slightly less customizable but has transparent PAYG pricing, structured flow agents built for bounded scripts, and faster time to a live pilot. For a PM-led pilot (not an engineering sprint), Retell wins. Trade-off: less customization for faster shipping and honest cost visibility.
AI identity	ALWAYS DISCLOSE AS AI	Some AI calling platforms allow the agent to avoid confirming it's AI. This would likely improve short-term connection rates. It was rejected immediately — Ditto's brand is built on no deception. If users later discovered they were talking to an undisclosed AI, the reputational damage would be irreversible. Trade-off: some calls will end at disclosure; brand safety is non-negotiable.
Funnel stages	FIRST-CONNECT + NO-SHOW ONLY	Document reminders and application follow-ups (Stage 6) are also automatable and would save real advisor time. They were deprioritized to the Phase 2 roadmap to keep the pilot scope narrow and the risk surface small. A narrower pilot is faster to evaluate, easier to kill if needed, and creates less exposure to brand damage. Trade-off: slower efficiency gains in Phase 1 for faster learning.
Script length	UNDER 60 SECONDS	The brief asked for under 60 seconds. More importantly, insurance leads are high-anxiety users who will disengage quickly if they sense a sales call. Shorter = more trust. Dia doesn't explain Ditto's history or list policy types — she does one thing: books a call. Trade-off: less information captured per call for a higher completion rate.
DNC handling	IMMEDIATE, PERMANENT STOP	Some call platforms allow a "re-engage in 30 days" logic after a DNC. This was rejected for Ditto. Once a user says "don't call me," the record is permanently flagged and no further AI contact is initiated. The economics of re-engaging a DNC lead do not justify the brand risk of being perceived as spam. Trade-off: lost lead recovery for stronger trust positioning.
Success metric	CSAT OVER BOOKINGS	A PM optimizing for consultation bookings alone could build a system that books lots of low-intent calls, wastes advisor time, and quietly erodes Ditto's 4.9 rating. The primary success test for this pilot is: did the 4.9 rating hold? Did CSAT after AI contact stay above 4.5? If yes, everything else is details. If no, the pilot failed regardless of booking volume.

If I Had More Time — 3 Specific Things I'd Do Differently

01

Live platform testing, not just pricing research

I'd get a trial account on both Retell and Vapi and make 20 test calls to Indian mobile numbers — specifically testing Hinglish mid-call switches, background noise handling, and latency. Pricing tables are inputs; a real call quality test is the output. The platform recommendation would be stronger with this data.

02

Interview one Ditto advisor about lead quality

The biggest risk to this pilot isn't the AI — it's Dia booking 100 consultations with low-intent leads who waste advisors' time. I'd spend 30 minutes with one advisor understanding what a "bad booked lead" looks like and how to detect low intent in the first 30 seconds of a call. That would sharpen the AI qualification script significantly.

03

A/B test two script variants in the pilot

The first-connect script I wrote is one approach. I'd run a second variant — shorter, with a different opener — to the same cohort of leads and compare booking rate and CSAT. Starting with a single script means the pilot takes longer to optimize. Two variants from day one halves the learning time.